

Experiment 4. Build a linear regression model using python on given data set by

- i. Prepare the data for ML model.
- ii. Splitting Training data and Test data.
- iii. Evaluate the model (intercept and slope).
- iv. Visualize the training set and testing set using Matplotlib, Seaborn.
- v. predicting the test set result.
- vi. compare actual output values with predicted values.

### Step 1: Import Libraries and Load the Dataset

We'll load the Boston Housing dataset, which is accessible online. We'll load it directly into a DataFrame using its URL.

```
# Import necessary libraries
import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error
import matplotlib.pyplot as plt
import seaborn as sns

# Load the dataset from a URL
url = "https://raw.githubusercontent.com/selva86/datasets/master/BostonHousing.csv"
data = pd.read_csv(url)

# Display the first few rows of the dataset
print(data.head())
```

```
➡
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	b	lstat	medv
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	5.33	36.2

## Step 2: Prepare Data for Modeling

The dataset contains various columns, with MEDV (Median value of owner-occupied homes) as the target variable. We'll use RM (average number of rooms per dwelling) as the input feature to predict MEDV.

```
# Step 2: Prepare data for ML model
# We are predicting the target variable 'medv' (median value of owner-occupied homes in $1000s)
# and using 'rm' (average number of rooms per dwelling) as the input feature.
X = data[['rm']]
Y = data['medv']
```

## Step 3: Train and Test Split

We'll split the data into training and testing sets for model training and evaluation

```
# Step 3: Splitting Training Data and Test Data
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, random_state=42)
```

## Step 4: Train and Evaluate the Model

Train a linear regression model on the training data and evaluate it by calculating the intercept and slope.

```
# Step 4: Train the model and evaluate the model (Intercept and Slope)
model = LinearRegression()
model.fit(X_train, Y_train)
```

```
# Model parameters
intercept = model.intercept_
slope = model.coef_[0]
```

```
print("Model Intercept:", intercept)
print("Model Slope:", slope)
```

```
➡ Model Intercept: -36.24631889813795
   Model Slope: 9.348301406497727
```

## Step 5: Visualize and Compare Results

We'll visualize both training and test predictions and compare actual values with predicted values.


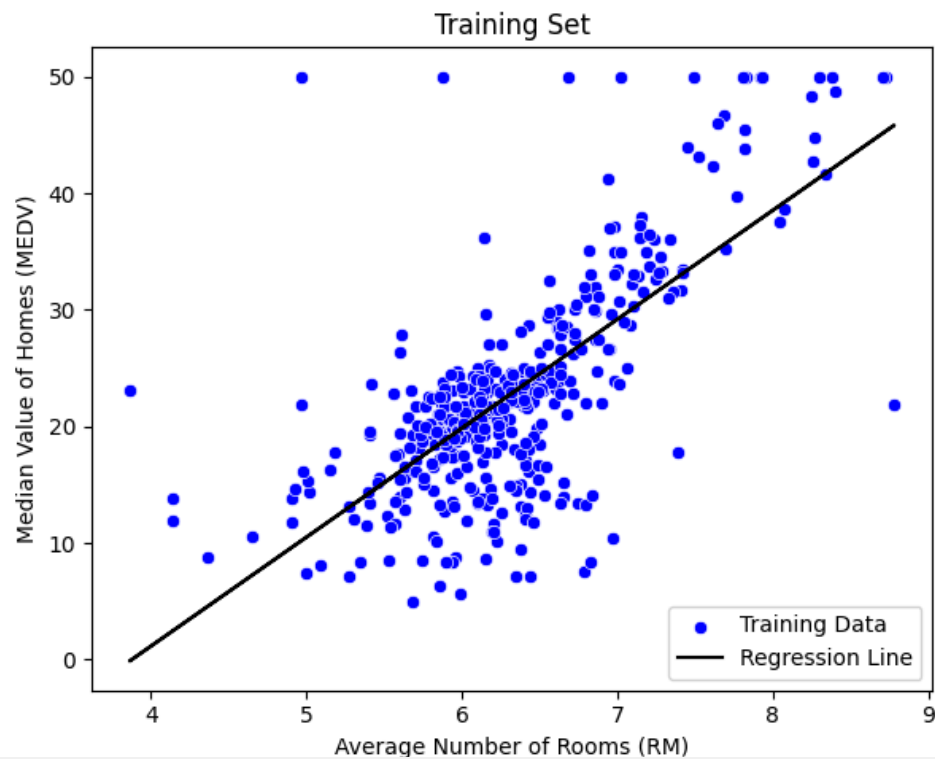
```
# Step 5: Visualize the Training Set and Testing Set
plt.figure(figsize=(12, 5))
```

```
# Training set visualization
```

```
plt.subplot(1, 2, 1)
plt.title("Training Set")
sns.scatterplot(x=X_train['rm'], y=Y_train, color='blue', label="Training Data")
plt.plot(X_train, model.predict(X_train), color='black', label="Regression Line")
plt.xlabel("Average Number of Rooms (RM)")
plt.ylabel("Median Value of Homes (MEDV)")
plt.legend()
```

```
# Testing set visualization
plt.subplot(1, 2, 2)
plt.title("Testing Set")
sns.scatterplot(x=X_test['rm'], y=Y_test, color='red', label="Testing Data")
plt.plot(X_test, model.predict(X_test), color='black', label="Regression Line")
plt.xlabel("Average Number of Rooms (RM)")
plt.ylabel("Median Value of Homes (MEDV)")
plt.legend()
```

```
plt.tight_layout()
plt.show()
```



```
# Predict the Test Set Results
Y_pred = model.predict(X_test)

# Compare Actual Output Values with Predicted Values
comparison = pd.DataFrame({'Actual': Y_test.values, 'Predicted': Y_pred})
print(comparison.head())

# Calculate model evaluation metrics
mse = mean_squared_error(Y_test, Y_pred)
mae = mean_absolute_error(Y_test, Y_pred)
print("Mean Squared Error:", mse)
print("Mean Absolute Error:", mae)
```

	Actual	Predicted
0	23.6	23.732383
1	32.4	26.929502
2	13.6	19.684568
3	22.8	20.451129
4	16.1	22.619935
Mean Squared Error:		46.144775347317264
Mean Absolute Error:		4.478335832064149